

머신비전시스템 과제 6

18011789 조혜수

1.

```
[40] import tensorflow as tf
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
```

```
[41] fashion_mnist = tf.keras.datasets.fashion_mnist
(train_X, train_Y), (test_X, test_Y) = fashion_mnist.load_data()
```

```
[42] train_X = train_X / 255.0
test_X = test_X / 255.0
```

(1)

1 : Trouser 4 : Coat 7 : Sneaker 8 : Bag

```
[43] index=[]
for i in range(len(train_Y)):
    if train_Y[i]==1 or train_Y[i]==4 or train_Y[i]==7 or train_Y[i]== 8:
        index.append(i)
```

```
y_train=[]
x_train=[]
```

```
for i in range(len(train_Y)):
    if i in index:
        x_train.append(train_X[i])
        y_train.append(train_Y[i])
```

```
▶ for i in range(len(y_train)):
    if y_train[i]==1:
        y_train[i]=0
    elif y_train[i]==4:
        y_train[i]=1
    elif y_train[i]==7:
        y_train[i]=2
    elif y_train[i]==8:
        y_train[i]=3
```

```
[45] print(len(x_train))
print(len(y_train))
```

```
24000
24000
```

```
[46] index=[]
    for i in range(len(test_Y)):
        if test_Y[i]==1 or test_Y[i]==4 or test_Y[i]==7 or test_Y[i]== 8:
            index.append(i)

    x_test=[]
    y_test=[]

    for i in range(len(test_Y)):
        if i in index:
            x_test.append(test_X[i])
            y_test.append(test_Y[i])
```

```
[47] for i in range(len(y_test)):
        if y_test[i]==1:
            y_test[i]=0
        elif y_test[i]==4:
            y_test[i]=1
        elif y_test[i]==7:
            y_test[i]=2
        elif y_test[i]==8:
            y_test[i]=3
```

```
[48] arr_train_x = np.array(x_train)
    arr_test_x = np.array(x_test)
```

(2)

```
[49] train_Yc = tf.keras.utils.to_categorical(y_train, num_classes=4)
     test_Yc = tf.keras.utils.to_categorical(y_test, num_classes=4)
```

```
[50] model = tf.keras.Sequential([
        tf.keras.layers.Flatten(input_shape=(28,28)),
        tf.keras.layers.Dense(units=128, activation='relu'),
        tf.keras.layers.Dense(units=64, activation='relu'),
        tf.keras.layers.Dense(units=32, activation='relu'),
        tf.keras.layers.Dense(units=4)
    ])
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
flatten_1 (Flatten)	(None, 784)	0
dense_4 (Dense)	(None, 128)	100480
dense_5 (Dense)	(None, 64)	8256
dense_6 (Dense)	(None, 32)	2080
dense_7 (Dense)	(None, 4)	132
=====		
Total params: 110,948		
Trainable params: 110,948		
Non-trainable params: 0		

(3)

```
[51] model.compile(optimizer=tf.keras.optimizers.SGD(),  
                  loss = 'mean_squared_error',  
                  metrics = ['accuracy'])
```

```
[52] history = model.fit(arr_train_x, train_Yc, batch_size=36, epochs=10)
```

```
Epoch 1/10  
667/667 [=====] - 2s 3ms/step - loss: 0.0504 - accuracy: 0.9277  
Epoch 2/10  
667/667 [=====] - 2s 3ms/step - loss: 0.0246 - accuracy: 0.9763  
Epoch 3/10  
667/667 [=====] - 2s 3ms/step - loss: 0.0203 - accuracy: 0.9813  
Epoch 4/10  
667/667 [=====] - 2s 3ms/step - loss: 0.0179 - accuracy: 0.9830  
Epoch 5/10  
667/667 [=====] - 2s 3ms/step - loss: 0.0163 - accuracy: 0.9838  
Epoch 6/10  
667/667 [=====] - 2s 3ms/step - loss: 0.0150 - accuracy: 0.9846  
Epoch 7/10  
667/667 [=====] - 2s 3ms/step - loss: 0.0140 - accuracy: 0.9850  
Epoch 8/10  
667/667 [=====] - 2s 3ms/step - loss: 0.0132 - accuracy: 0.9859  
Epoch 9/10  
667/667 [=====] - 2s 3ms/step - loss: 0.0125 - accuracy: 0.9863  
Epoch 10/10  
667/667 [=====] - 2s 3ms/step - loss: 0.0119 - accuracy: 0.9869
```

```
[53] model.evaluate(arr_test_x, test_Yc)
```


```
125/125 [=====] - 0s 2ms/step - loss: 0.0125 - accuracy: 0.9870  
[0.012451119720935822, 0.9869999885559082]
```

(4) Trouser, Coat, Sneaker, Bag

✓
1초

```
[54] #Trouser
      img = arr_test_x[62]
      pred = model.predict(tf.expand_dims(img, axis=0))
      plt.imshow(img, 'gray')
      print(y_test[62], pred)


0 [[ 0.99963176  0.05685138 -0.03616047  0.03575461]]
```



✓
0초

```
#coat
      img = arr_test_x[500]
      pred = model.predict(tf.expand_dims(img, axis=0))
      plt.imshow(img, 'gray')
      print(y_test[500], pred)


☞ 1 [[ 0.07180193  0.83958006 -0.08399118  0.11062928]]
```



```
[56] #Sneaker
```

```
img = arr_test_x[70]  
pred = model.predict(tf.expand_dims(img, axis=0))  
plt.imshow(img, 'gray')  
print(y_test[70], pred)
```


```
2 [[-4.16612960e-02  2.16025859e-04  9.07469869e-01  1.08501315e-01]]
```



```
#bag
```

```
img = arr_test_x[550]  
pred = model.predict(tf.expand_dims(img, axis=0))  
plt.imshow(img, 'gray')  
print(y_test[550], pred)
```

```
3 [[0.06645723 0.0546686 0.01786293 0.9252195 ]]
```



2.

```
[20] import tensorflow as tf
import matplotlib.pyplot as plt
import numpy as np
```



```
fashion_mnist = tf.keras.datasets.fashion_mnist
(train_X, train_Y), (test_X, test_Y) = fashion_mnist.load_data()
```

```
[22] train_X = train_X / 255.0
test_X = test_X / 255.0
```

(1)

```
[23] index=[]
for i in range(len(train_Y)):
    if train_Y[i]==1 or train_Y[i]==4 or train_Y[i]==7 or train_Y[i]== 8:
        index.append(i)
```

```
y_train=[]
x_train=[]
```

```
for i in range(len(train_Y)):
    if i in index:
        x_train.append(train_X[i])
        y_train.append(train_Y[i])
```

```
[24] for i in range(len(y_train)):
    if y_train[i]==1:
        y_train[i]=0
    elif y_train[i]==4:
        y_train[i]=1
    elif y_train[i]==7:
        y_train[i]=2
    elif y_train[i]==8:
        y_train[i]=3
```

```
[25] index=[]
    for i in range(len(test_Y)):
        if test_Y[i]==1 or test_Y[i]==4 or test_Y[i]==7 or test_Y[i]== 8:
            index.append(i)

    x_test=[]
    y_test=[]

    for i in range(len(test_Y)):
        if i in index:
            x_test.append(test_X[i])
            y_test.append(test_Y[i])

[26] for i in range(len(y_test)):
    if y_test[i]==1:
        y_test[i]=0
    elif y_test[i]==4:
        y_test[i]=1
    elif y_test[i]==7:
        y_test[i]=2
    elif y_test[i]==8:
        y_test[i]=3

[27] arr_train_x = np.array(x_train)
    arr_train_y = np.array(y_train)
    arr_test_x = np.array(x_test)
    arr_test_y = np.array(y_test)
```

(2)

```
[28] model = tf.keras.Sequential([
    tf.keras.layers.Conv2D(input_shape=(28,28,1), kernel_size=(5,5),
        strides=(1, 1), padding='same', filters=20),
    tf.keras.layers.Activation('relu'),
    tf.keras.layers.MaxPooling2D(pool_size=(2,2), strides=(2,2), padding='valid'),
    tf.keras.layers.Conv2D(kernel_size=(5,5),
        strides=(1, 1), padding='same', filters=50), tf.keras.layers.Activation('relu'),
    tf.keras.layers.MaxPooling2D(pool_size=(2,2), strides=(2,2), padding='valid'), tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(units=500),
    tf.keras.layers.Activation('relu'),
    tf.keras.layers.Dense(units=4),
    tf.keras.layers.Softmax()
])
```


(3)

```
[29] model.compile(optimizer=tf.keras.optimizers.Adam(),
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])

[30] history = model.fit(arr_train_x, arr_train_y, batch_size=36, epochs=10, validation_split=0.25)

Epoch 1/10
500/500 [=====] - 43s 84ms/step - loss: 0.0694 - accuracy: 0.9789 - val_loss: 0.0368 - val_accuracy: 0.9880
Epoch 2/10
500/500 [=====] - 42s 84ms/step - loss: 0.0260 - accuracy: 0.9926 - val_loss: 0.0275 - val_accuracy: 0.9915
Epoch 3/10
500/500 [=====] - 44s 88ms/step - loss: 0.0187 - accuracy: 0.9942 - val_loss: 0.0221 - val_accuracy: 0.9935
Epoch 4/10
500/500 [=====] - 41s 82ms/step - loss: 0.0115 - accuracy: 0.9963 - val_loss: 0.0296 - val_accuracy: 0.9923
Epoch 5/10
500/500 [=====] - 41s 83ms/step - loss: 0.0086 - accuracy: 0.9975 - val_loss: 0.0233 - val_accuracy: 0.9950
Epoch 6/10
500/500 [=====] - 43s 86ms/step - loss: 0.0068 - accuracy: 0.9981 - val_loss: 0.0204 - val_accuracy: 0.9948
Epoch 7/10
500/500 [=====] - 41s 83ms/step - loss: 0.0054 - accuracy: 0.9983 - val_loss: 0.0315 - val_accuracy: 0.9937
Epoch 8/10
500/500 [=====] - 42s 84ms/step - loss: 0.0039 - accuracy: 0.9986 - val_loss: 0.0326 - val_accuracy: 0.9917
Epoch 9/10
500/500 [=====] - 43s 85ms/step - loss: 0.0029 - accuracy: 0.9989 - val_loss: 0.0359 - val_accuracy: 0.9943
Epoch 10/10
500/500 [=====] - 41s 83ms/step - loss: 0.0041 - accuracy: 0.9988 - val_loss: 0.0419 - val_accuracy: 0.9923


[31] model.evaluate(arr_test_x, arr_test_y, verbose=False)

[0.025588931515812874, 0.9937499761581421]
```

(4)


```
[32] img = arr_test_x[62]
     pre_img = tf.expand_dims(img, axis=0)
     pre_img = pre_img/255
     pred = model.predict(pre_img)
     plt.imshow(img, 'gray')
     np.argmax(pred)
```

2



```
[33] img = arr_test_x[500]
     pre_img = tf.expand_dims(img, axis=0)
     pre_img = pre_img/255
     pred = model.predict(pre_img)
     plt.imshow(img, 'gray')
     np.argmax(pred)
```


2



```
img = arr_test_x[70]  
pre_img = tf.expand_dims(img, axis=0)  
pre_img = pre_img/255  
pred = model.predict(pre_img)  
plt.imshow(img, 'gray')  
np.argmax(pred)
```




2



```
img = arr_test_x[550]  
pre_img = tf.expand_dims(img, axis=0)  
pre_img = pre_img/255  
pred = model.predict(pre_img)  
plt.imshow(img, 'gray')  
np.argmax(pred)
```



2



3.

```
[2] import tensorflow as tf
import numpy as np
from google.colab.patches import cv_imshow
import matplotlib.pyplot as plt
```

```
[3] (train_X, train_Y), (test_X, test_Y) = tf.keras.datasets.cifar10.load_data()
print(train_X.shape, test_X.shape)
```

```
Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170500096/170498071 [=====] - 3s 0us/step
170508288/170498071 [=====] - 3s 0us/step
(50000, 32, 32, 3) (10000, 32, 32, 3)
```

```
[4] train_X = tf.keras.applications.vgg16.preprocess_input(train_X)
test_X = tf.keras.applications.vgg16.preprocess_input(test_X)
```

(1)

```
[5] index=[]
    for i in range(len(train_Y)):
        if train_Y[i]==0 or train_Y[i]==1 or train_Y[i]==8:
            index.append(i)

    y_train=[]
    x_train=[]

    for i in range(len(train_Y)):
        if i in index:
            x_train.append(train_X[i])
            y_train.append(train_Y[i])
```

```
[6] for i in range(len(y_train)):
        if y_train[i]==0:
            y_train[i]=0
        elif y_train[i]==1:
            y_train[i]=1
        elif y_train[i]==8:
            y_train[i]=8
```

```
[7] index=[]
    for i in range(len(test_Y)):
        if test_Y[i]==0 or test_Y[i]==1 or test_Y[i]==8:
            index.append(i)

    x_test=[]
    y_test=[]

    for i in range(len(test_Y)):
        if i in index:
            x_test.append(test_X[i])
            y_test.append(test_Y[i])
```

```
[8] for i in range(len(y_test)):
        if y_test[i]==0:
            y_test[i]=0
        elif y_test[i]==1:
            y_test[i]=1
        elif y_test[i]==8:
            y_test[i]=8
```

(2)

```
▶ model = tf.keras.applications.VGG16(include_top=False)
model.summary()
```

```
↳ Downloading data from https://storage.googleapis.com/tensorflow/keras-applic:
58892288/58889256 [=====] - 0s 0us/step
58900480/58889256 [=====] - 0s 0us/step
Model: "vgg16"
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, None, None, 3)]	0
block1_conv1 (Conv2D)	(None, None, None, 64)	1792
block1_conv2 (Conv2D)	(None, None, None, 64)	36928
block1_pool (MaxPooling2D)	(None, None, None, 64)	0
block2_conv1 (Conv2D)	(None, None, None, 128)	73856
block2_conv2 (Conv2D)	(None, None, None, 128)	147584
block5_conv1 (Conv2D)	(None, None, None, 512)	2359808
block5_conv2 (Conv2D)	(None, None, None, 512)	2359808
block5_conv3 (Conv2D)	(None, None, None, 512)	2359808
block5_pool (MaxPooling2D)	(None, None, None, 512)	0

=====
Total params: 14,714,688
Trainable params: 14,714,688
Non-trainable params: 0

```
[10] base_model = tf.keras.applications.VGG16(input_shape=[32,32,3],
                                              include_top=False, weights='imagenet')

x = base_model.output
x = tf.keras.layers.Flatten()(x)
x = tf.keras.layers.Dense(64, activation='relu')(x)
predictions = tf.keras.layers.Dense(10, activation='softmax')(x)
model = tf.keras.Model(inputs=base_model.input, outputs=predictions)
```

(3)

```
▶ for layer in model.layers[:19]:  
    layer.trainable = False  
for layer in model.layers[19:]:  
    layer.trainable = True  
model.summary()
```

☞ Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_2 (InputLayer)	[(None, 32, 32, 3)]	0
block1_conv1 (Conv2D)	(None, 32, 32, 64)	1792
block1_conv2 (Conv2D)	(None, 32, 32, 64)	36928
block1_pool (MaxPooling2D)	(None, 16, 16, 64)	0
block2_conv1 (Conv2D)	(None, 16, 16, 128)	73856
block2_conv2 (Conv2D)	(None, 16, 16, 128)	147584
block2_pool (MaxPooling2D)	(None, 8, 8, 128)	0
block3_conv1 (Conv2D)	(None, 8, 8, 256)	295168
block3_conv2 (Conv2D)	(None, 8, 8, 256)	590080
block3_conv3 (Conv2D)	(None, 8, 8, 256)	590080
block3_pool (MaxPooling2D)	(None, 4, 4, 256)	0
dense (Dense)	(None, 64)	32832
dense_1 (Dense)	(None, 10)	650
=====		
Total params: 14,748,170		
Trainable params: 33,482		
Non-trainable params: 14,714,688		

```
[12] model.compile(optimizer=tf.keras.optimizers.Adam(),  
                  loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

```
▶ arr_train_x = np.array(x_train)  
arr_train_y = np.array(y_train)  
arr_test_x = np.array(x_test)  
arr_test_y = np.array(y_test)
```

```
[16] history = model.fit(arr_train_x, arr_train_y, batch_size=36, epochs=10, validation_split=0.25)
```

```
Epoch 1/10  
313/313 [=====] - 16s 16ms/step - loss: 1.4074 - accuracy: 0.7565 - val_loss: 0.6971 - val_accuracy: 0.8099  
Epoch 2/10  
313/313 [=====] - 4s 14ms/step - loss: 0.4673 - accuracy: 0.8501 - val_loss: 0.5602 - val_accuracy: 0.8184  
Epoch 3/10  
313/313 [=====] - 4s 14ms/step - loss: 0.3105 - accuracy: 0.8834 - val_loss: 0.5563 - val_accuracy: 0.8301  
Epoch 4/10  
313/313 [=====] - 4s 14ms/step - loss: 0.2335 - accuracy: 0.9129 - val_loss: 0.5513 - val_accuracy: 0.8427  
Epoch 5/10  
313/313 [=====] - 4s 14ms/step - loss: 0.1912 - accuracy: 0.9282 - val_loss: 0.5566 - val_accuracy: 0.8405  
Epoch 6/10  
313/313 [=====] - 4s 14ms/step - loss: 0.1542 - accuracy: 0.9436 - val_loss: 0.6057 - val_accuracy: 0.8363  
Epoch 7/10  
313/313 [=====] - 5s 15ms/step - loss: 0.1273 - accuracy: 0.9541 - val_loss: 0.6375 - val_accuracy: 0.8373  
Epoch 8/10  
313/313 [=====] - 4s 14ms/step - loss: 0.1006 - accuracy: 0.9637 - val_loss: 0.6917 - val_accuracy: 0.8392  
Epoch 9/10  
313/313 [=====] - 4s 14ms/step - loss: 0.0931 - accuracy: 0.9681 - val_loss: 0.7338 - val_accuracy: 0.8347  
Epoch 10/10  
313/313 [=====] - 4s 14ms/step - loss: 0.0856 - accuracy: 0.9696 - val_loss: 0.8142 - val_accuracy: 0.8347
```

```
[18] model.evaluate(arr_test_x, arr_test_y, verbose=False)
```

```
[0.7662954330444336, 0.8320000171661377]
```


(4)

```
[11] for layer in model.layers:
      layer.trainable = True
      model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 32, 32, 3)]	0
block1_conv1 (Conv2D)	(None, 32, 32, 64)	1792
block1_conv2 (Conv2D)	(None, 32, 32, 64)	36928

```
=====
Total params: 14,748,170
Trainable params: 14,748,170
Non-trainable params: 0
=====
```

```
[12] model.compile(optimizer=tf.keras.optimizers.Adam(),
                  loss='sparse_categorical_crossentropy', metrics=['accuracy'])
      history = model.fit(arr_train_x, arr_train_y, batch_size=36, epochs=10, validation_split=0.25)
```

```
Epoch 1/10
313/313 [=====] - 24s 35ms/step - loss: 1.4447 - accuracy: 0.4043 - val_loss: 0.9792 - val_accuracy: 0.4896
Epoch 2/10
313/313 [=====] - 10s 32ms/step - loss: 0.7797 - accuracy: 0.5970 - val_loss: 0.5919 - val_accuracy: 0.7192
Epoch 3/10
313/313 [=====] - 10s 32ms/step - loss: 0.5339 - accuracy: 0.7766 - val_loss: 0.3978 - val_accuracy: 0.8336
Epoch 4/10
313/313 [=====] - 10s 33ms/step - loss: 0.3567 - accuracy: 0.8550 - val_loss: 0.2689 - val_accuracy: 0.8965
Epoch 5/10
313/313 [=====] - 10s 33ms/step - loss: 0.2721 - accuracy: 0.8964 - val_loss: 0.2953 - val_accuracy: 0.8936
Epoch 6/10
313/313 [=====] - 10s 33ms/step - loss: 0.2092 - accuracy: 0.9253 - val_loss: 0.3961 - val_accuracy: 0.8813
Epoch 7/10
313/313 [=====] - 10s 33ms/step - loss: 0.1834 - accuracy: 0.9332 - val_loss: 0.2504 - val_accuracy: 0.9168
Epoch 8/10
313/313 [=====] - 10s 33ms/step - loss: 0.1652 - accuracy: 0.9436 - val_loss: 0.2038 - val_accuracy: 0.9251
Epoch 9/10
313/313 [=====] - 10s 33ms/step - loss: 0.1237 - accuracy: 0.9576 - val_loss: 0.2006 - val_accuracy: 0.9301
Epoch 10/10
313/313 [=====] - 10s 33ms/step - loss: 0.1651 - accuracy: 0.9439 - val_loss: 0.2089 - val_accuracy: 0.9304
```

```
[13] model.evaluate(arr_test_x, arr_test_y, verbose=False)
```

```
[0.2004273682832718, 0.9279999732971191]
```